

Applied Geo-Scripting

Jan Verbesselt, et al.

October 3, 2013

1 Geo-scripting learning objectives

- Handle spatial data using a scripting language (e.g. R)
- Understand basic concepts of applied scripting for spatial data
- Learn to read, write, and visualize spatial data (vector/raster) using a script
- Know how to find help (on spatial data handling functions)
- Solve scripting problems (debug, reproducible example, writing functions)
- Find libraries which offer spatial data handling functions
- Learn to include functions from a library in your script
- Apply learned concepts in a case study: learning how to address a spatial/ecological/applied case (e.g. detect forest changes, flood mapping, ocean floor depth analysis, bear movement, etc.) with a raster and vector dataset.

2 Today's Learning objectives

- Understand basic concepts of applied scripting for spatial data
- Handle spatial data using a scripting language (e.g. R)

3 Why geo-scripting?

- Reproducible: avoid clicking and you keep track of what you have done
- Efficient: you can write a script to do something for you e.g. multiple times e.g. automatically downloading data
- Enable collaboration: sharing scripts, functions, and packages
- Good for finding errors i.e. debugging
e.g. this course is fully written with scripting languages (i.e. R and LaTeX).

4 What is a scripting language?

A scripting language or script language is a programming language that supports the writing of scripts, programs written for a special runtime environment that can interpret and automate the execution of tasks which could alternatively be executed one-by-one by a human operator. Different from compiled languages like C/C++/Fortran.

A scripting language is the glue, between different commands, functions, and objectives without the need to compile it for each OS/CPU Architecture.

5 Different scripting languages for geo-scripting

The main scripting languages for GIS and Remote sensing currently are: R, Python (stand-alone or integrated within ArcGIS, QGIS), GRASS.

Sytze, Aldo, ... can you add more info here

6 Python versus R

- Python is a general purpose programming language
- R is particularly strong in statistical computing and graphics
- Installing libraries in Python is sometimes challenging
- Syntactic differences can be confusing
- There are many R packages for spatial analyses and for dealing with spatial data
- Scripts in both languages can be combined

Sytze can you help here...

7 Course set-up and planning

- R package: RASTA package https://r-forge.r-project.org/R/?group_id=1743
- Have look at the Reproducible and Applied Spatial and Temporal Analysis (RASTA) package (package content)
- Course set-up is that every lesson there will be a short introduction, followed by a tutorial and an exercise that needs to be handed in before the start of the next lesson.
- Course content and overview

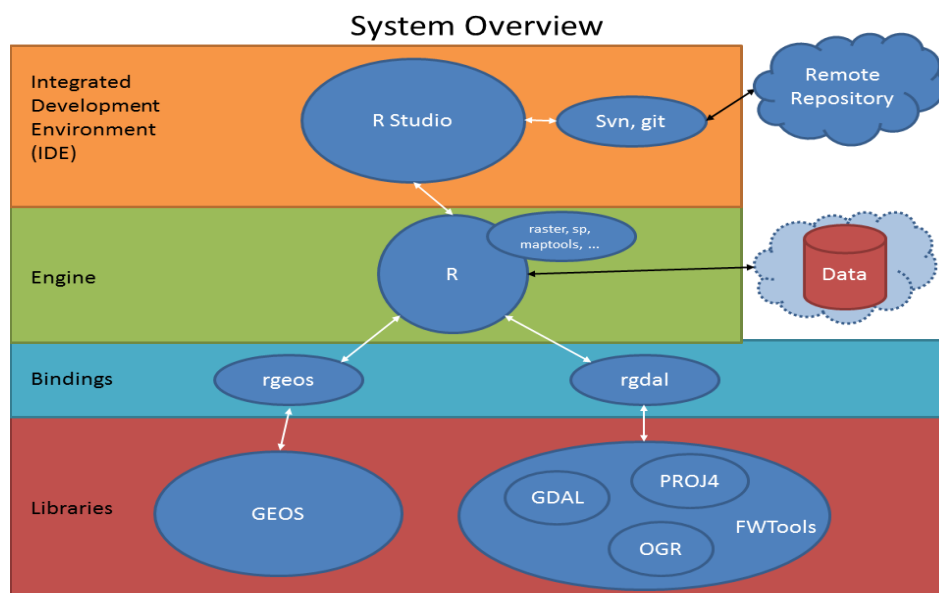


Figure 1: The graphical user interface to R

8 Get Your R On

8.1 Hello world with Rstudio

This preliminary section will cover some basic details about R. For this course we will use Rstudio as an IDE to write and run scripts.

When you open Rstudio you will see Fig. x The window in the bottom left corner is the R console (e.g. statistical and spatial analysis tools). Now type the following script in the R console:

```
> a <- 1
> a

[1] 1

> print(Hello World)

[1] "Hello World"
```

The first line you passed to the console created a new object named *a* in memory. The symbol '`<-`' is somewhat equivalent to an equal sign. In the second line you printed *a* to the console by simply typing its name.

Now try to obtain the following output in the R console by writing the commands in the script window and running them via **Ctrl-r** (make sure you remove `>` and `+`). R commands shown in the following section should be written in your Script file without `>` and `+` in front of it, so that you can run them by using **Ctrl-r** and the result in the R console will look like this:

```
> class(a)

[1] "numeric"

> class(Hello World)

[1] "character"
```

You now have requested the **class** attribute of *a* and the console has returned the attribute: **numeric**. R possesses a simple mechanism to support an object-oriented style of programming. All objects (*a* in this case) have a class attribute assigned to them. **R** is quite forgiving and will assign a class to an object even if you haven't specified one (as you didn't in this case). Classes are a very important feature of the **R** environment. Any function or method that is applied to an object takes into account its class and uses this information to determine the correct course of action. A simple example should suffice to explain further.

Select the next two lines using your mouse and pass these to the console using **Ctrl-r**. The first line passed declares a new object *b*. The second line passed adds *a* and *b* together and prints the solution to the console.

```
> b <- 2
> a + b

[1] 3
```

R has assessed the class attribute of *a* and *b*; determined they are both **numeric** objects, and; carried out the arithmetic calculation as requested.

8.2 Custom Functions

It is hard to unleash the full potential of R without writing your own functions. Luckily it's very easy to do. Here are some trivial examples:

```
> #-----CUSTOM FUNCTIONS-----
>
> add<-function(x){ #put the function arguments in () and the evaluation in {}
+   x+1
+ }
> add(3)

[1] 4

> add(4)

[1] 5

> #--Set the default values for your function--
> add<-function(x=5){
+   x+1
+ }
> add() #automatically evaluates x=5

[1] 6

> add(6) #but you can still change the defaults

[1] 7
```

That's about all there is too it. The function will generally return the result of the last line that was evaluated. However you can also use `return()` to specify exactly what the function will return.

Now, let's declares a new object, a new function, **newfunc** (this is just a name and if you like you can give this function another name). Appearing in the first set of brackets is an argument list that specifies (in this case) two names. The value of the function appears within the second set of brackets where the process applied to the named objects from the argument list is defined.

```
> newfunc <- function(x, y) {
+   2*x + y
+ }
> a2b <- newfunc(2, 4)
> a2b

[1] 8

> rm(a, b, newfunc, a2b)
```

Next, a new object `a2b` is created which contains the result of applying **newfunc** to the two objects you have defined earlier. The second last R command prints this new object to the console. Finally, you can now remove the objects you have created to make room for the next exercise by selecting and running the last line of the code.

8.3 Help?!

R is supported by a very comprehensive help system. Help on any function can be accessed by entering the name of the function into the console preceded with a ?. The easiest way to access the system is to open a web-browser. This help system can be started by entering `help.start()` in the R console. Try it and see what happens.

```
> # ?class
```

8.4 Data Structures

There are several ways that data are stored in R. Here are the main ones:

- **Vectors** The most generic data structure. In R, any variable of an atomic data type (numeric, integer, logical, character) is a vector. This will be exemplified below.
- **Data Frames** The most common format. Similar to a spread sheet. A `data.frame()` is indexed by rows and columns and store numeric and character data. The `data.frame` is typically what we use when we read in csv files, do regressions, et cetera.
- **Matrices and Arrays** Similar to `data.frames` but slightly faster computation wise while sacrificing some of the flexibility in terms of what information can be stored. In R a matrix object is a special case of an array that only has 2 dimensions. IE, an array is n-dimensional matrix while a matrix only has rows and columns (2 dimensions)
- **Lists** The most common and flexible type of R object. A list is simply a collection of other objects. For example a regression object is a list of: 1) Coefficient estimates 2) Standard Errors 3) The Variance/Covariance matrix 4) The design matrix (data) 5) Various measures of fit, et cetera.

We will look at examples of these objects in the next section

8.5 Reading Data in and Out

The most common way to read in data is with the `read.csv()` command. However you can read in virtually any type of text file. Type `?read.table` in your console for some examples. If you have really large binary data sets sometimes the `scan()` function is more efficient. Finally using the foreign package you can read in SPSS, STATA, Matlab, SAS, and a host of other data formats from other stat and math software.

Let's read in a basic csv file.

```
> #-----READING DATA IN AND OUT-----  
> library(rasta) ## load the rasta library  
> f <- system.file("extdata/kenpop89to99.csv", package="rasta")  
> ## make a link to the csv file within the rasta package  
> mydat<-read.csv(f)
```

We can explore the data using the `names()`, `summary()`, `head()`, and `tail()` commands (we will use these frequently through out the exercise)

```
> names(mydat) #column names
```

```
[1] "ip89DId"      "ip89DName"  "ADMIN3"      "KEADMN3_ID" "Y89Pop"
[6] "Y89Births"    "Y89Brate"   "Y99Pop"      "Y99Births"   "Y99Brate"
[11] "PopChg"       "BrateChg"
```

```
> summary(mydat) #basic summary information
```

ip89DId		ip89DName		ADMIN3	KEADMN3_ID
Min.	:1010	Kisii	: 3	KISII	: 2
1st Qu.:	3772	Kakamega	: 2	BARINGO	: 1
Median	:6010	Kericho	: 2	BOMET	: 1
Mean	:5207	Machakos	: 2	BUNGOMA	: 1
3rd Qu.:	7052	Meru	: 2	BUSIA	: 1
Max.	:8030	South Nyanza:	2	E. MARAKWET:	1
		(Other)	:35	(Other)	:41

Y89Pop		Y89Births		Y89Brate		Y99Pop	
Min.	: 57960	Min.	: 1680	Min.	:22.64	Min.	: 72380
1st Qu.:	222905	1st Qu.:	9350	1st Qu.:	33.52	1st Qu.:	392545
Median	: 451510	Median	:18270	Median	:37.38	Median	: 629740
Mean	: 619710	Mean	:23719	Mean	:37.03	Mean	: 872928
3rd Qu.:	947500	3rd Qu.:	39855	3rd Qu.:	40.88	3rd Qu.:	1384665
Max.	:1476500	Max.	:57460	Max.	:51.01	Max.	:2363120

Y99Births		Y99Brate		PopChg		BrateChg	
Min.	: 1760	Min.	:19.03	Min.	: -14.00	Min.	: -38.00
1st Qu.:	10870	1st Qu.:	28.03	1st Qu.:	23.75	1st Qu.:	-20.00
Median	:21820	Median	:31.01	Median	: 33.50	Median	: -14.00
Mean	:27562	Mean	:31.57	Mean	: 47.73	Mean	: -14.56
3rd Qu.:	42140	3rd Qu.:	36.36	3rd Qu.:	44.25	3rd Qu.:	-6.75
Max.	:69380	Max.	:42.89	Max.	:343.00	Max.	: 0.00

```
> head(mydat) #first 6 rows
```

	ip89DId	ip89DName	ADMIN3	KEADMN3_ID	Y89Pop	Y89Births	Y89Brate	Y99Pop
1	1010	Nairobi	NAIROBI	41	1325620	42560	32.11	2085820
2	2010	Kiambu	KIAMBU	38	908120	27720	30.52	1383300
3	2020	Kirinyaga	KIRINYAGA	29	389440	10980	28.19	452180
4	2030	Muranga	MURANGA	36	862540	27940	32.39	737520
5	2040	Nyandaura	NYANDARUA	22	348520	12520	35.92	468300
6	2050	Nyeri	NYERI	26	607980	17540	28.85	644380

	Y99Births	Y99Brate	PopChg	BrateChg
1	58700	28.14	57	-12
2	36140	26.13	52	-14
3	10840	23.97	16	-15
4	16500	22.37	-14	-31
5	13320	28.44	34	-21
6	14340	22.25	6	-23

```
> tail(mydat) # last 6 rows
```

	ip89DId	ip89DName	ADMIN3	KEADMN3_ID	Y89Pop	Y89Births	Y89Brate	
43	7120	Uasin-Gishu	UASIN	GISHU	13	443280	17900	40.38
44	7130	West-Pokot	WEST	POKOT	5	224640	9440	42.02
45	8010	Bugoma	BUNGOMA		11	741940	34600	46.63
46	8020	Busia	BUSIA		16	425380	18640	43.82
47	8030	Kakamega	VIHIGA		21	1476500	57460	38.92
48	8030	Kakamega	KAKAMEGA		14	1476500	57460	38.92

	Y99Pop	Y99Births	Y99Brate	PopChg	BrateChg
43	616240	22260	36.12	39	-11
44	309020	12940	41.87	38	0
45	1008080	43240	42.89	36	-8
46	547680	23440	42.80	29	-2
47	2011960	69380	34.48	36	-11
48	2011960	69380	34.48	36	-11

We will go over ways to index and subscript data.frames later on in the exercise. For now lets do a basic regression so you can see an example of a list

8.6 Basic Regression (and an example of lists)

We use the `lm()` command to do a basic linear regression. The `~` symbol separates the left and right hand sides of the equation and we use `+` to separate terms and `*` to specify interactions.

```
> #-----REGRESSION AND LISTS-----
> myreg<-lm(Y99Pop~Y89Births+Y89Brate,data=mydat) #Regress the Population in 1999 on the p
> myreg
```

Call:

```
lm(formula = Y99Pop ~ Y89Births + Y89Brate, data = mydat)
```

Coefficients:

(Intercept)	Y89Births	Y89Brate
502592.59	38.05	-14369.09

A regression object is an example of a list. We can use the `names()` command to see what the list contains. We can use the `summary()` command to get a standard regression output (coefficients, standard errors, et cetera) and we can also create a new object that contains all the elements of a regression summary.

```
> #-----EXPLORE A REGRESSION OBJECT-----
>
> names(myreg) #get the names of the items in the regression object

[1] "coefficients" "residuals"      "effects"         "rank"
[5] "fitted.values" "assign"          "qr"              "df.residual"
[9] "xlevels"      "call"           "terms"           "model"

> summary(myreg) #print out the key information
```



```

Call:
lm(formula = Y99Pop ~ Y89Births + Y89Brate, data = mydat)

Residuals:
    Min       1Q   Median       3Q      Max
-362649 -117800 -10240   36497  597511

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) 502592.593 199219.406   2.523  0.0152 *
Y89Births      38.045     2.028  18.762 <2e-16 ***
Y89Brate    -14369.093   5774.647  -2.488  0.0166 *
---
Signif. codes:  0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

Residual standard error: 215200 on 45 degrees of freedom
Multiple R-squared: 0.8983,    Adjusted R-squared: 0.8938
F-statistic: 198.7 on 2 and 45 DF,  p-value: < 2.2e-16

> myregsum<-summary(myreg) #create a new regression summary object
> names(myregsum)

[1] "call"          "terms"          "residuals"      "coefficients"
[5] "aliased"        "sigma"          "df"             "r.squared"
[9] "adj.r.squared" "fstatistic"     "cov.unscaled"

> myregsum[[adj.r.squared]] #extract the adjusted r squared

[1] 0.8937546

> myregsum$adj.r.squared #does the same thing

[1] 0.8937546

```

That concludes our basic introduction to data.frames and lists. There is alot more material out on the web if you are interested. Later in the exercise we will look at data.frames in more detail.

9 Excercise 1: Write you own function to do something fancy

10 More information

For more information about R please refer to the following links <http://www.statmethods.net/index.html>. This is a great website for learning R function, graphs, and stats. Also visit <http://www.r-project.org/> and check out the Manuals i.e an introductions to R See also the book on Applied spatial Data analysis with R <http://www.asdar-book.org/> (Bivand et al., 2013).

References

Bivand, R. S., Pebesma, E. J., & Rubio, V. G. (2013). Applied spatial data analysis with R, .