

Raster Analysis

Ben DeVries, Jan Verbesselt, Loïc Dutrieux, Sytze de Bruin

October 16, 2013

Abstract

In this tutorial, we will explore the raster package and other related packages used for typical raster analyses. We will first look at analysis of RasterLayer objects, exploring functions having to do with raster algebra, focal and zonal statistics and other operations. We will then explore spatio-temporal analysis of raster data using RasterBrick objects. Here, we will extract time series data from a RasterBrick object and derive temporal statistics from these data. Since data from the Landsat archive is becoming increasingly important for environmental research and monitoring, this tutorial will focus on the use of these data.

1. perform raster algebra to calculate indices
2. classify a raster layer
3. perform focal operations to sieve a raster
4. parse Landsat scene information from time series data
5. explore a raster brick by plotting layers and layers statistics
6. perform raster brick operations to derive statistics (e.g. %no-data in the time series)
7. extract pixel time series and derive various time series statistics

1 Raster and related packages

The raster package is an essential tool for raster-based analysis in R. Here you will find functions which are fundamental to image analysis. The raster package documentation is a good place to begin exploring the possibilities of image analysis within R. There is also an excellent vignette available at <http://cran.r-project.org/web/packages/raster/vignettes/Raster.pdf>.

In addition to the raster package, we will be using the rasterVis and ggplot2 packages to make enhanced plots.

```
> # load the necessary packages
> library(raster)
> library(rasterVis)
> library(ggplot2)
```

2 The Landsat archive

Since being released to the public, the Landsat data archive has become an invaluable tool for environmental monitoring. With a historical archive reaching back to the 1970's, the release of these data has resulted in a spur of time series based methods. In this tutorial, we will work with time series data from the Landsat 7 Enhanced Thematic Mapper (ETM+) sensor.

3 Manipulating raster data

Exploring a Landsat scene

Landsat scenes are delivered via the USGS as a number of image layers representing the different bands captured by the sensors. In the case of the Landsat 7 Enhanced Thematic Mapper (ETM+) sensor, the bands are shown in figure xxx. Using different combination of these bands can be useful in describing land features and change processes.

TODO: insert figure showing Landsat bands.

Part of a landsat scene, including bands 2-4 are included as part of the rasta package. These data have been processed using the LEDAPS framework (TODO: insert link), so the values contained in this dataset represent surface reflectance, scaled by 10000 (ie. divide by 10000 to get a reflectance value between 0 and 1).

We will begin exploring these data simply by visualizing them (more methods for data exploration will be covered in Lesson 7).

```
> # load in the data
> data(GewataB2)
> data(GewataB3)
> data(GewataB4)
> # check out the attributes
> GewataB2

class      : RasterLayer
dimensions : 1177, 1548, 1821996  (nrow, ncol, ncell)
resolution : 30, 30  (x, y)
extent     : 808755, 855195, 817635, 852945  (xmin, xmax, ymin, ymax)
coord. ref.: +proj=utm +zone=36 +ellps=WGS84 +units=m +no_defs
data source: in memory
names      : gewataB2
values     : 127, 2220  (min, max)

> GewataB3

class      : RasterLayer
dimensions : 1177, 1548, 1821996  (nrow, ncol, ncell)
resolution : 30, 30  (x, y)
extent     : 808755, 855195, 817635, 852945  (xmin, xmax, ymin, ymax)
coord. ref.: +proj=utm +zone=36 +ellps=WGS84 +units=m +no_defs
data source: in memory
names      : gewataB3
values     : 85, 2865  (min, max)
```

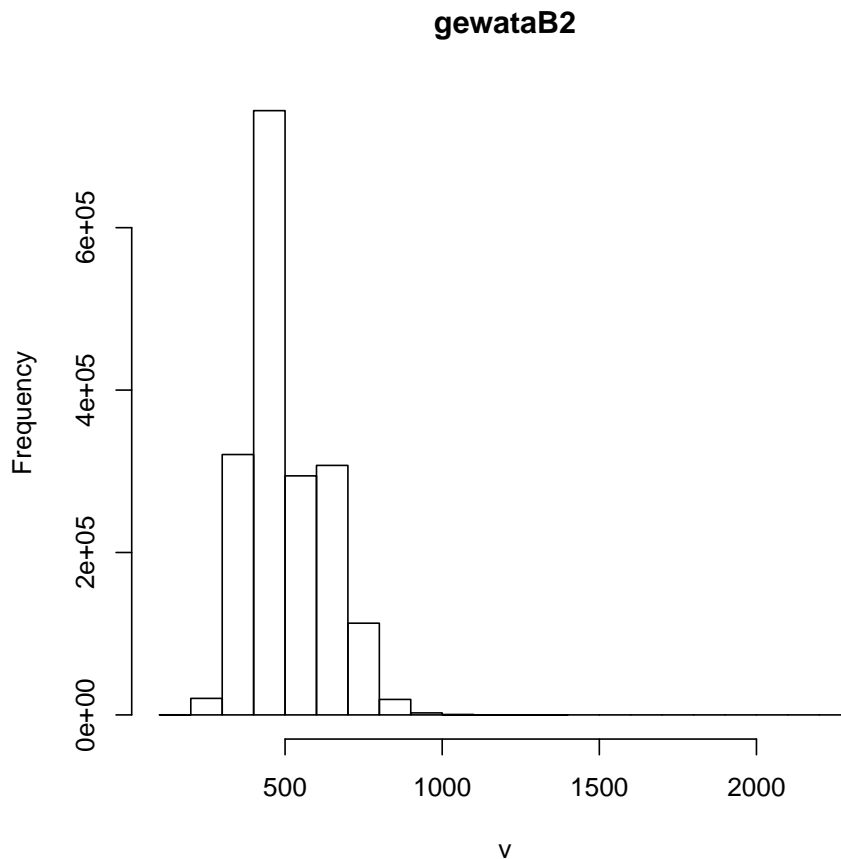
```

> GewataB4

class       : RasterLayer
dimensions  : 1177, 1548, 1821996  (nrow, ncol, ncell)
resolution  : 30, 30  (x, y)
extent      : 808755, 855195, 817635, 852945  (xmin, xmax, ymin, ymax)
coord. ref. : +proj=utm +zone=36 +ellps=WGS84 +units=m +no_defs
data source : in memory
names       : gewataB4
values      : 598, 4903  (min, max)

> # plot the histograms of all bands
> hist(GewataB2)
> hist(GewataB3)
> hist(GewataB4)

```

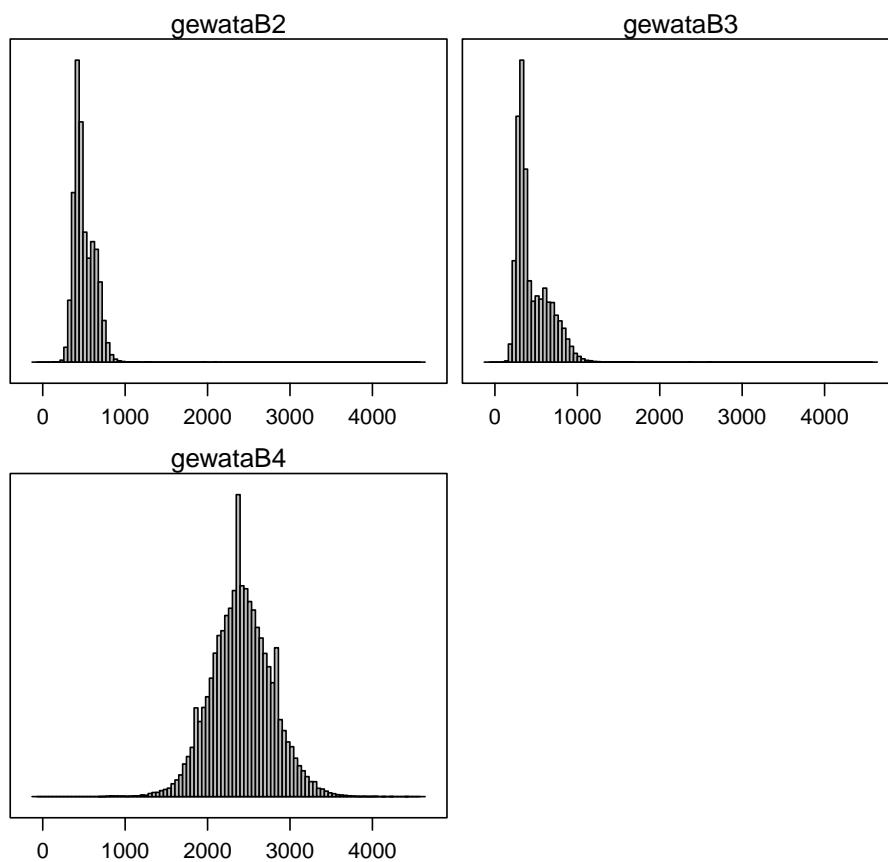


We can improve these plots by adjusted the scale, etc., but the rasterVis package that has enhanced plotting capabilities which make it easier to make more attractive plots. First, to make the comparison easier, we will make a rasterBrick object from these three layers.

```

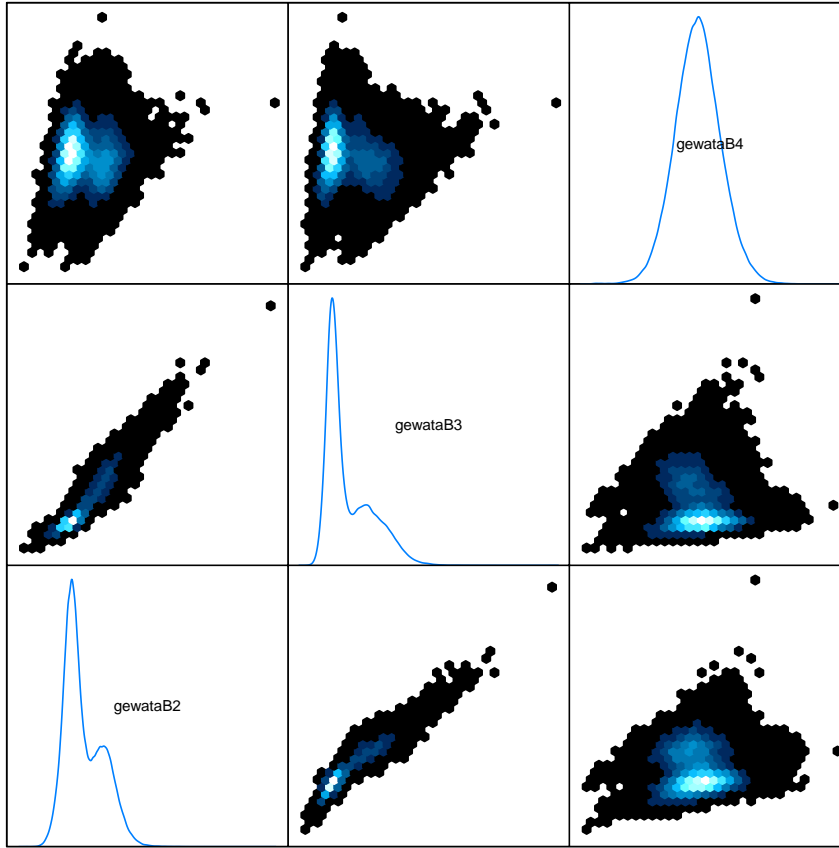
> gewata <- brick(GewataB2, GewataB3, GewataB4)
> # view all histograms together with rasterVis
> histogram(gewata)

```



The rasterVis package has several other raster plotting types inherited from the lattice package. For multispectral data, one plot type which is particularly useful for data exploration is the scatterplot matrix, called by the `splom()` function.

```
> splom(gewata)
```



Scatter Plot Matrix

Calling `splom()` on a `rasterBrick` reveals potential correlations between the layers themselves. In the case of bands 2-4 of the `gewata` subset, we can see that band 2 and 3 (in the visual part of the EM spectrum) are highly correlated, while band 4 contains significant non-redundant information. This band, situated in the near infrared (NIR) region of the EM spectrum, is often used to describe vegetation-related features for this reason.

The `rasterVis` package will be demonstrated in more detail in this tutorial the section dealing with multilayered raster data.

Computing new rasters: Raster algebra

In the previous section, we observed a tight correlation between two of the Landsat bands of the `gewata` subset, but a very different distribution of values in band 4 (NIR). This distribution stems from the fact that vegetation reflects very highly in the NIR range, compared to the visual range of the EM spectrum. A commonly used metric for assessing vegetation dynamics, the normalized difference vegetation index (NDVI), takes advantage of this fact and is computed from Landsat bands 3 (red; R) and 4 (NIR) as follows:

$$NDVI = \frac{NIR - R}{NIR + R} \quad (1)$$

Using principles of raster algebra, we can easily perform this calculation in R.

```
> ndvi <- (GewataB4 - GewataB3) / (GewataB4 + GewataB3)
```

Using raster algebra (above) is relatively easy and intuitive. However, with particularly large datasets or complex calculations, it is often desirable to use the `overlay()` function instead.

```
> ndvi <- overlay(GewataB4, GewataB3, fun = function(x,y){(x-y)/(x+y)})
```

One advantage of this function is the fact that the result can be written immediately to file by including the argument `'filename=...'`, thus saving memory (especially important when working with large datasets).

Plotting our new raster, we can immediately see how useful the NDVI metric is for identifying land features. Use the interactive `'drawExtent()'` function to zoom into some of the features to inspect them more closely.

```
> # first, plot the raster
> plot(ndvi)
> # call drawExtent() to activate interactive mode
> # and assign the result to an extent object e
> e <- drawExtent()
> # now click 2 points on the plot window
> # these represent the top-right and bottom-left corner of your extent
> # now plot ndvi again, but only the extent you defined interactively
> plot(ndvi, ext = e)
```

4 Classifying raster data

One of the most important tasks in analysis of remote sensing image analysis is image classification. In classifying the image, we take the information contained in the various bands (possibly including other synthetic bands such as NDVI or principle components). In this tutorial we will explore two approaches for image classification: unsupervised (k-means) and supervised (random forest) classification.

Supervised classification: Random Forest

...TODO: add description of the random forest classifier, with links...

We should first prepare the data on which the classification will be done. So far, we have prepared three bands from a ETM+ image in 2001 (bands 2, 3 and 4) as a `rasterBrick`, and have also calculated NDVI. In addition, there is a Vegetation Continuous Field (VCF) product available for the same period (2000). This product is also based on Landsat ETM+ data, and represents an estimate of tree cover (in %). Since this layer could also be useful in classifying land cover types, we will also include it as a potential covariate in the random forest classification.

```
> # load the data
> data(vcfGewata)
> plot(vcfGewata)
> histogram(vcfGewata) # or 'hist(vcfGewata)'
```

```

> # note that there are some strange values >> 100
> # let's replace these values with NA
> vcfGewata[vcfGewata > 100] <- NA
> plot(vcfGewata)
> histogram(vcfGewata)

```

To perform the classification in R, it is best to assemble all covariate layers into one rasterBrick object. In this case, we can simply append these new layers (NDVI and VCF) to our existing rasterBrick (currently consisting of bands 2, 3, and 4). But first, let's rescale the NDVI layer by 10000 (just as the others have been scaled) and store it as an integer raster.

```

> ndvi <- calc(ndvi, fun = function(x) floor(x*10000))
> # change the data type
> # see ?dataType for more info
> dataType(ndvi) <- "INT2U"
> names(ndvi) <- "NDVI"
> # make the covariate rasterBrick
> covs <- addLayer(gewata, ndvi, vcfGewata)
> plot(covs)

```

For this exercise, we will do a very simple classification for 2001 using three classes: forest, cropland and wetland. While for other purposes it is usually better to define more classes (and possibly fuse classes later), a simple classification like this one could be useful, for example, to construct a forest mask for the year 2001.

```

> # load the training polygons
> data(trainingPoly)
> # inspect the data
> trainingPoly@data

```

	OBJECTID	Class
0	1	wetland
1	2	wetland
2	3	wetland
3	4	wetland
4	5	wetland
5	6	forest
6	7	forest
7	8	forest
8	9	forest
9	10	forest
10	11	cropland
11	12	cropland
12	13	cropland
13	14	cropland
14	15	cropland
15	16	cropland

```
> plot(ndvi)
> plot(trainingPoly, add = TRUE)
```

The training classes are labelled as string labels. For this exercise, we will need to work with integer classes, so we will need to first 'relabel' our training classes.

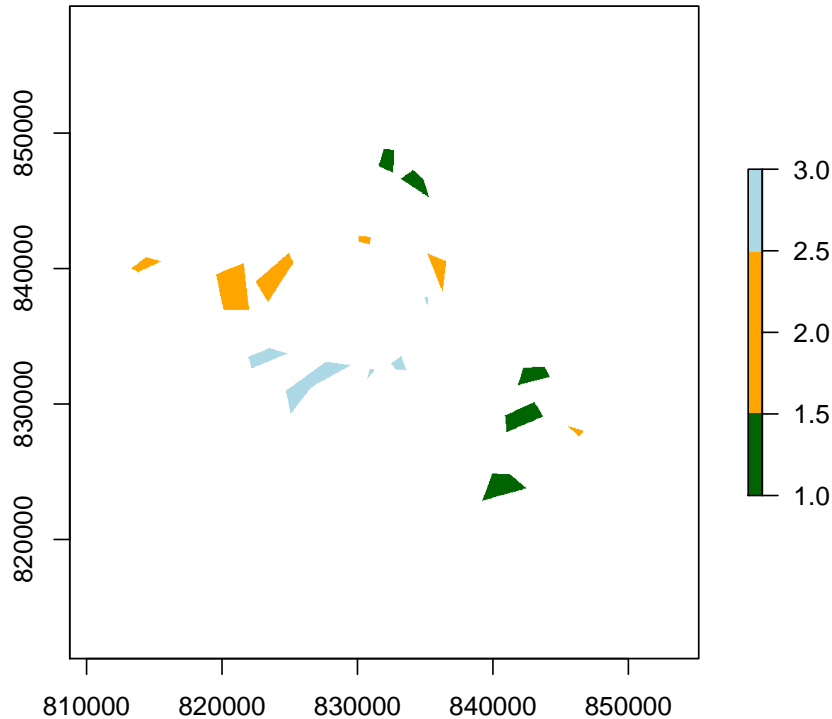
```
> # define a reclassification function
> reclass <- function(x){
+   y <- ifelse(x == "forest", 1,
+               ifelse(x == "cropland", 2, 3)
+               )
+   return(y)
+ }
> # apply this over the trainingPoly data slot
> trainingPoly@data$Code <- sapply(trainingPoly@data$Class, FUN = function(x) reclass(x))
```

To train the raster data, we need to convert our training data to the same type using the `rasterize()` function. This function takes a spatial object (in this case a polygon object) and transfers the values to raster cells defined by a raster object. Here, we will define a new raster containing those values.

```
> classes <- rasterize(trainingPoly, ndvi, field = 'Code')
```

Found 16 region(s) and 16 polygon(s)

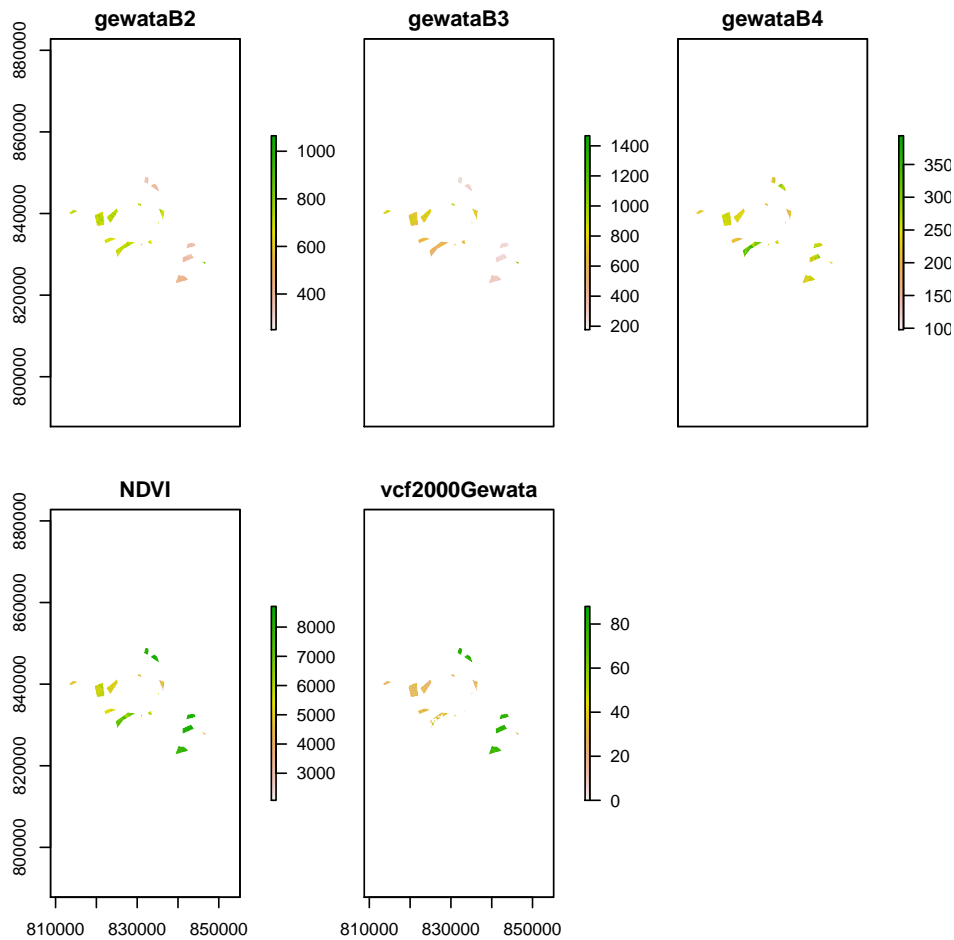
```
> dataType(classes) <- "INT1U"
> plot(classes, col = c("dark green", "orange", "light blue"))
```

(Note: there is a handy 'progress = "text"' argument, which can be passed to many of the raster package functions and can help to monitor processing).

Our goal in preprocessing these data is to have a table of values representing all layers (covariates) with *known* values/classes. To do this, we will first need to create a version of our rasterBrick only representing the training pixels. Here the mask() function from the raster package will be very useful.

```
> covmasked <- mask(covs, classes)
> plot(covmasked)
> # add the classes layer to this new brick
> names(classes) <- "class"
> trainingbrick <- addLayer(covmasked, classes)
> plot(trainingbrick)
```



Now it's time to add all of these values to a data.frame representing all training data. This data.frame will be used as an input into the RandomForest classification function. We will use `getValues()` to extract all of the values from the layers of the rasterBrick.

```
> # extract all values into a matrix
> valuetable <- getValues(trainingbrick)
> # convert to a data.frame
> valuetable <- as.data.frame(valuetable)
> head(valuetable)
```

	gewataB2	gewataB3	gewataB4	NDVI	vcf2000Gewata	class
1	NA	NA	NA	NA	NA	NA
2	NA	NA	NA	NA	NA	NA
3	NA	NA	NA	NA	NA	NA
4	NA	NA	NA	NA	NA	NA
5	NA	NA	NA	NA	NA	NA
6	NA	NA	NA	NA	NA	NA

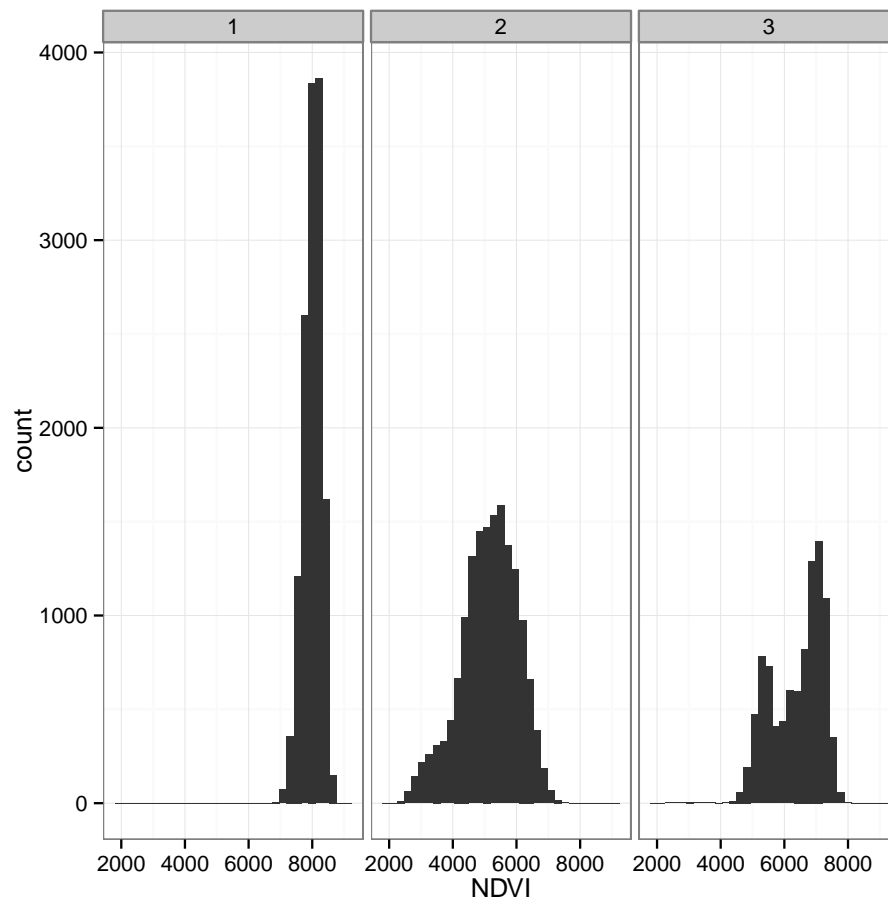
```
> # keep only rows where valuetable$class has a value
> valuetable <- valuetable[!is.na(valuetable$class),]
> head(valuetable)
```

	gewataB2	gewataB3	gewataB4	NDVI	vcf2000Gewata	class
212849	354	281	2527	7998	77	1
212850	417	301	2816	8068	74	1
212851	418	282	2857	8203	72	1
212852	397	282	2651	8077	71	1
212853	355	263	2237	7896	77	1
214397	396	281	2693	8110	75	1

```
> # convert values in the class column to factors
> valuetable$class <- factor(valuetable$class, levels = c(1:3))
```

Now we have a convenient reference table which contains, for each of the three defined classes, all known values for all covariates. Let's visualize the distribution of some of these covariates using the `ggplot()` package.

```
> # 1. NDVI
> ggplot(data = valuetable, aes(x = NDVI)) +
+   geom_histogram() +
+   facet_wrap(~ class) +
+   theme_bw()
> # 2. VCF
> ggplot(data = valuetable, aes(x = vcf2000Gewata)) +
+   geom_histogram() +
+   labs(x = "% Tree Cover") +
+   facet_wrap(~ class) +
+   theme_bw()
> # 3. Bands 3 and 4
> ggplot(data = valuetable, aes(x = gewataB3, y = gewataB4)) +
+   stat_bin2d() +
+   facet_wrap(~ class) +
+   theme_bw()
> # 4. Bands 2 and 3
> ggplot(data = valuetable, aes(x = gewataB2, y = gewataB3)) +
+   stat_bin2d() +
+   facet_wrap(~ class) +
+   theme_bw()
```



We can see from these distributions that these covariates may do well in classifying forest pixels, but we may expect some confusion between cropland and wetland (although the individual bands may help to separate these classes). When performing this classification on large datasets and with a large amount of training data, now may be a good time to save this table using the `write.csv()` command, in case something goes wrong after this point and you need to start over again.

Now it is time to build the Random Forest model using the training data contained in the table of values we just made. For this, we will use the "randomForest" package in R, which is an excellent resource for building such types of models. Using the `randomForest()` function, we will build a model based on a matrix of predictors or covariates (ie. the first 5 columns of `valuetable`) related to the response (the 'class' column of `valuetable`).

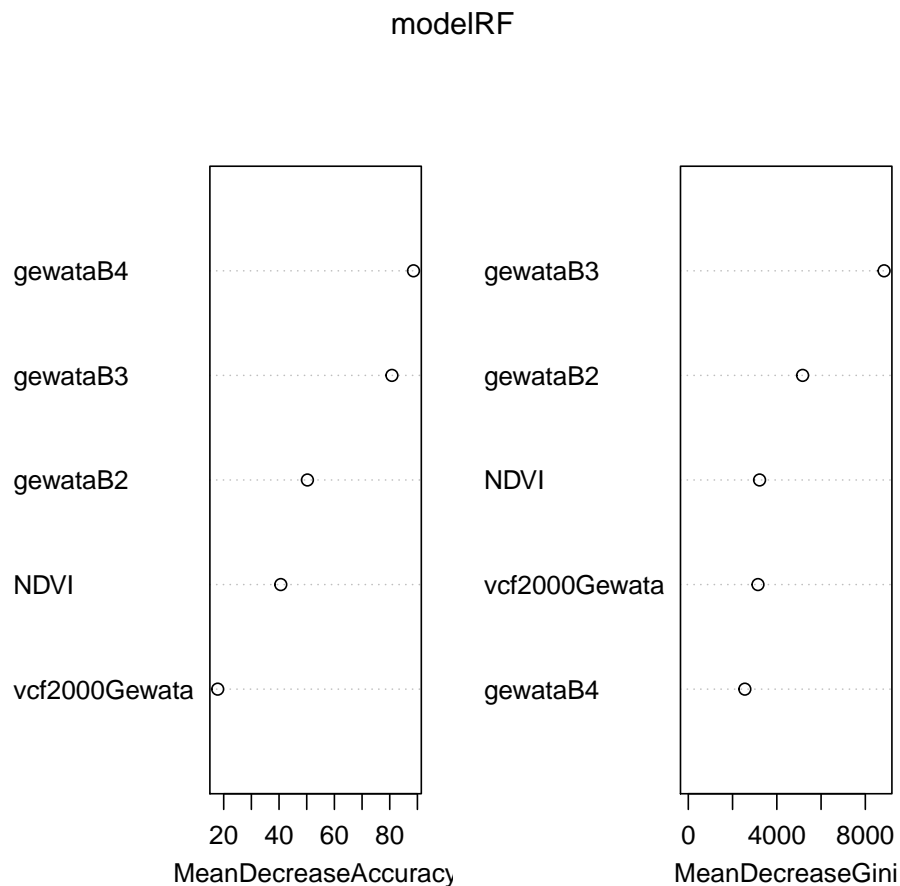
```
> # NA values are not permitted in the covariates/predictor columns
> # which rows have NAs in them?
> delRows <- which(apply(valuetable, 1, FUN = function(x) NA %in% x))
> # remove these rows from valuetable
> valuetable <- valuetable[-delRows,]
> # construct a random forest model
> # caution: this step takes fairly long!
> library(randomForest)
> modelRF <- randomForest(x = valuetable[,c(1:5)], y = valuetable$class,
```

```
+ importance = TRUE)
```

Since the random forest method involves the building and testing of many classification trees (the 'forest'), it is a computationally expensive step (and could take a lot of memory for especially large training datasets). When this step is finished, it would be a good idea to save the resulting object with the `save()` command. Any R object can be saved as an `.rda` file and reloaded into future sessions.

The resulting object from the `randomForest()` function is quite large and packed full of information. Since we set `'importance = TRUE'`, we now also have information on the statistical importance of each of our covariates which we can visualize using the `varImpPlot()` command.

```
> varImpPlot(modelRF)
```



...TODO: clarify this plot and show the OOB accuracy assessment...

Now we can apply this model to the rest of the image and assign classes to all pixels. Note that for this step, the names of the raster layers in the input brick (here 'covs') must correspond to the column names of the training table. We will use the `predict()` function from the `raster` package to predict class values based on the random forest model we have just constructed. This function uses a pre-defined model to predict values of raster cells based on other raster layers. This model can be derived by a linear regression, for example. In our case, we will use the model provided by the `randomForest()` function we applied earlier.

```

> # check layer and column names
> names(covs)

[1] "gewataB2"      "gewataB3"      "gewataB4"      "NDVI"
[5] "vcf2000Gewata"

> names(valuetable)

[1] "gewataB2"      "gewataB3"      "gewataB4"      "NDVI"
[5] "vcf2000Gewata" "class"

> # predict land cover using the RF model
> predLC <- predict(covs, model = modelRF, na.rm = TRUE)
> # plot the results
> # recall: 1 = forest, 2 = cropland, 3 = wetland
> plot(predLC, col = c("dark green", "orange", "light blue"))

```

Note that the `predict()` function also takes arguments that can be passed to `writeRaster()` (eg. `'filename = ...'`), so it would be a good idea to write to file as you perform this step (rather than keeping all output in memory).

Unsupervised classification: k-means

...paragraph on k-means...

We will use the same layers (from the 'covs' rasterBrick) as in the Random Forest classification for this classification exercise. As before, we need to extract all values into a data.frame.

```

> valuetable <- getValues(covs)
> head(valuetable)

```

	gewataB2	gewataB3	gewataB4	NDVI	vcf2000Gewata
[1,]	519	520	2479	6532	49
[2,]	563	577	2310	6002	37
[3,]	585	674	2140	5209	28
[4,]	563	694	2054	4949	28
[5,]	651	771	2096	4621	26
[6,]	739	943	2351	4274	19

Now we will construct a `kmeans` object using the `kmeans()` function. Like the Random Forest model, this object packages useful information about the resulting class membership. In this case, we will set the number of clusters to three, presumably corresponding to the three classes defined in our random forest classification.

```

> km <- kmeans(na.omit(valuetable), centers = 3, iter.max = 100, nstart = 10)
> # km contains the clusters (classes) assigned to the cells
> head(km$cluster)

[1] 1 1 2 2 2 2

```

```
> unique(km$cluster) # displays unique values
```

```
[1] 1 2 3
```

Here, we used the 'na.omit()' argument to avoid any NA values in the valuetable (recall that there is a region of NAs in the VCF layer). These NAs are problematic in the kmeans() function, but omitting them gives us another problem: the resulting vector of clusters (from 1 to 3) is shorter than the actual number of cells in the raster. In other words: how do we know which clusters to assign to which cells? To answer that question, we need to have a kind of 'mask' raster, indicating where the NA values throughout the cov rasterBrick are located.

```
> # create a blank raster with NA values
> rNA <- setValues(raster(covs), NA)
> # loop through layers of covs
> # assign a 1 wherever an NA is encountered
> for(i in 1:nlayers(covs)){
+   rNA[is.na(covs[[i]])] <- 1
+ }
> # convert rNA to an integer vector
> rNA <- getValues(rNA)
> # substitute the NA's for 0's
> rNA[is.na(rNA)] <- 0
```

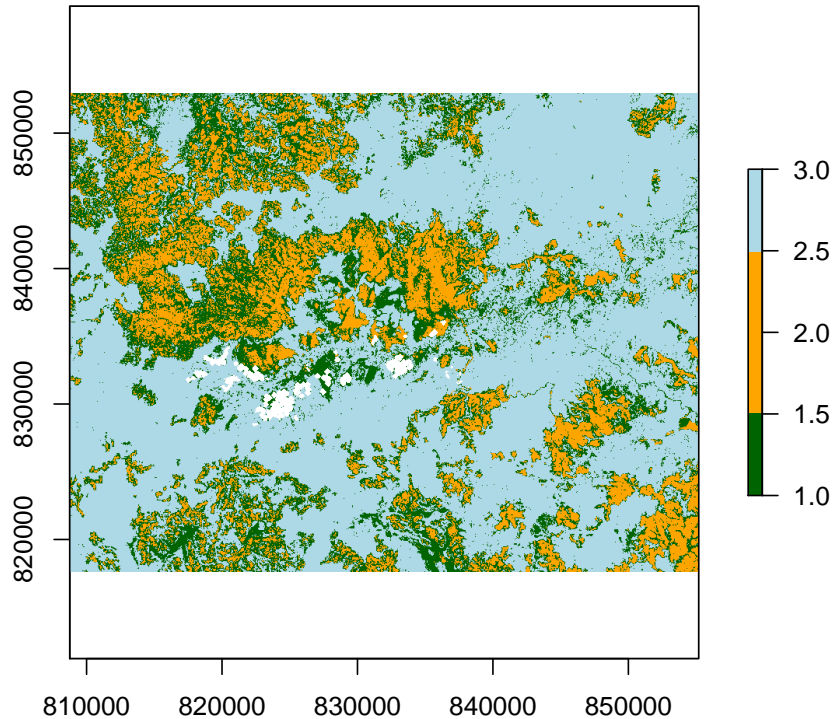
We now have a vector indicating with a value of 1 where the NA's in the cov brick are. Now that we know where the 'original' NAs are located, we can go ahead and assign the cluster values to a raster. At these 'NA' locations, we will not assign any of the cluster values, instead assigning an NA.

First, we will insert these values into the original valuetable data.frame.

```
> # convert valuetable to a data.frame
> valuetable <- as.data.frame(valuetable)
> # assign the cluster values (where rNA != 1)
> valuetable$class[rNA==0] <- km$cluster
> # assign NA to this column elsewhere
> valuetable$class[rNA==1] <- NA
```

Now we are finally ready to assign these cluster values to a raster. This will represent our final classified raster.

```
> # create a blank raster
> classes <- raster(covs)
> # assign values from the 'class' column of valuetable
> classes <- setValues(classes, valuetable$class)
> plot(classes, col = c("dark green", "orange", "light blue"))
```



These classes are much more difficult to interpret than those resulting from the random forest classification. Clearly, with a good training dataset, a supervised classification can provide a reasonably accurate land cover classification. However, unsupervised classification methods like k-means are useful for study areas for which little to no *a priori* data exist. Is there a way we could improve the k-means classification performed in this example? Which one is computationally faster between random forest and k-means (hint: try the `system.time()` function)?

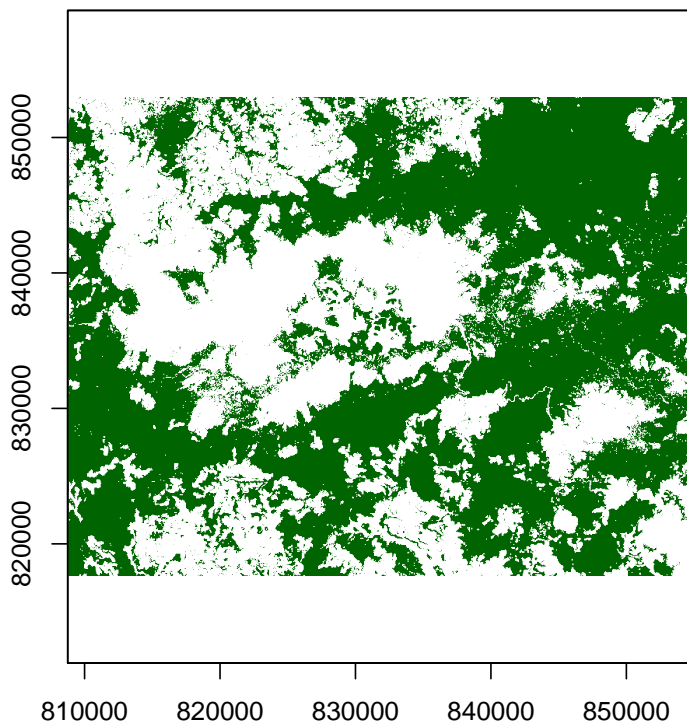
Applying a raster sieve

Although the land cover raster we created above is limited in the number of thematic classes it has, and we observed some confusion between wetland and cropland classes, it could be useful for constructing a forest mask. To do so, we have to fuse (and remove) non-forest classes, and then clean up the remaining pixels using focal algebra by applying a sieve.

```
> # Make an empty raster based on the LC raster attributes
> formask <- raster(predLC)
> # assign NA to all cells
> formask <- setValues(formask, value = NA)
> # assign 1 to all cells corresponding to the forest class
> formask[predLC == 1] <- 1
```



```
> plot(formask, col = "dark green", legend = FALSE)
```



We now have a forest mask that can be used to isolate forest pixels for further analysis. For some applications, however, we may only be interested in larger forest areas. We may especially want to remove single forest pixels, as they may be a result of errors, or may not fit our definition of 'forest'.

To remove these pixel 'islands', we will use the `focal()` function of the raster package, which computes values based on values of neighbouring pixels. The first task is to construct a mask with the `focal()` function, which will be used to 'clean' up the forest mask we have just produced.

In the first case, we will define 'island' pixels as any pixel not having neighbours in any direction (including diagonals). To do this, we can simply apply a single numeric weight defining the dimensions of the matrix.

```
> # make an empty raster
> sievemask <- setValues(raster(formask), NA)
> # assign a value of 1 for all forest pixels
> sievemask[!is.na(formask)] <- 1
> # sum of all neighbourhood pixels
```

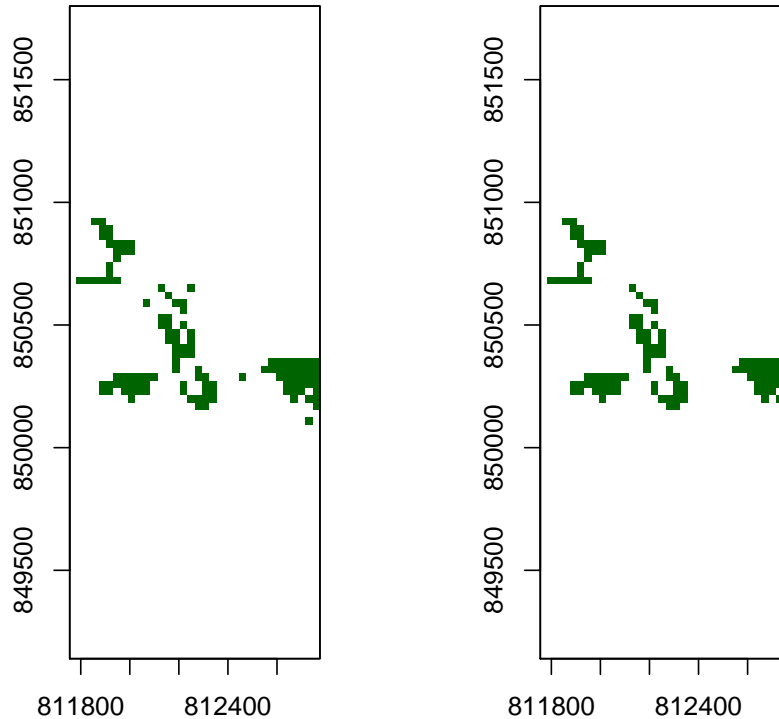
```
> sievemask <- focal(siemask, w = 3, fun = sum, na.rm = TRUE)
> sievemask
```

```
class      : RasterLayer
dimensions : 1177, 1548, 1821996  (nrow, ncol, ncell)
resolution : 30, 30  (x, y)
extent     : 808755, 855195, 817635, 852945  (xmin, xmax, ymin, ymax)
coord. ref.: +proj=utm +zone=36 +ellps=WGS84 +units=m +no_defs
data source : in memory
names      : layer
values     : 0, 1  (min, max)
```

```
> histogram(siemask)
```

We now have a mask whose values are the sum of neighbourhood weights (each equal to $1/9$, based on a 3×3 window). In cases where there are no neighbours (ie. the 'island' pixels), this sum will be equal to $1/9$ exactly. In cases with one or more neighbours, the sum will be greater than $1/9$, up to 1 (ie. if the pixel is completely surrounded by non-NA pixels). To apply the sieve and remove 'island' pixels, we want to select for only cases where `sievmask == 1/9` and remove those from the original raster.

```
> # copy the original forest mask
> formaskSieve <- formask
> # assign NA to pixels where the sievmask == 1/9
> formaskSieve[sievmask == 1/9] <- NA
> # zoom in to a small extent to check the results
> # Note: you can define your own by using e <- drawExtent()
> e <- extent(c(811744.8, 812764.3, 849997.8, 850920.3))
> par(mfrow = c(1, 2)) # allow 2 plots side-by-side
> plot(formask, ext = e, col = "dark green", legend = FALSE)
> plot(formaskSieve, ext = e, col = "dark green", legend = FALSE)
> par(mfrow = c(1, 1)) # reset plotting window
```



We have successfully removed all 'island' pixels from the forest mask using the `focal()` function. Suppose we define 'island' pixels as those have no immediate neighbours, *not* considering diagonal neighbours. In that case, we would have to adjust the weight argument ('w') in `focal()`, and instead define our own 3X3 matrix which omits diagonal pixels.

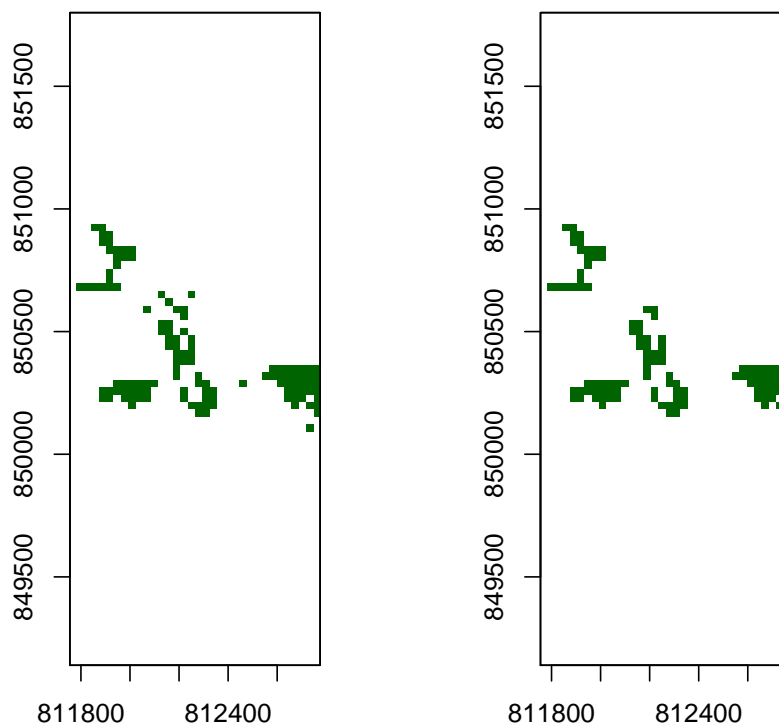
```
> # define a weights matrix
> w <- rbind(c(0, 1, 0),
+           c(1, 1, 1),
+           c(0, 1, 0))
> print(w)

      [,1] [,2] [,3]
[1,]    0    1    0
[2,]    1    1    1
[3,]    0    1    0

> # alternatively:
> w <- matrix(c(0, 1, 0, 1, 1, 1, 0, 1, 0), nrow = 3)
> # prepare the sievemask (as above)
> sievemask <- setValues(raster(formask), NA)
> sievemask[!is.na(formask)] <- 1
```

```
> # sum of all neighbouring pixels, except for diagonals
> sievemask <- focal(siemask, w = w, fun = sum, na.rm = TRUE)
```

When we define the matrix manually, the weights are not normalized (sum to 1) as they did when we set 'w = 3'. In this case, 'island' pixels would have a value of 1 (no other neighbours), and should be removed according to the sieve. Applying the new sieve as above should give the following results.



In this case, not only pixels with absolutely no neighbours have been removed, but also those pixels with only diagonal neighbours as well. The second case could be considered as a more 'conservative' sieve.

The `focal()` function can be used for a variety of purposes, including other data filters such as median, Laplacian, Sobel, etc. More complex weight matrices can also be computed, such as a Gaussian using the `focalWeight()` function. These methods will not be covered in this tutorial. Check out the documentation at `?focal()` for more information.

Working with thematic rasters

In some cases, the values of a raster may be categorical, meaning they relate to a thematic class (e.g. 'forest' or 'wetland') rather than a quantitative value (e.g. NDVI or % Tree Cover). The raster dataset 'lulcGewata' is a raster with integer values representing LULC classes from a 2011 classification (using SPOT5 and ASTER source data).

```
> data(lulcGewata)
> # check out the distribution of the values
> freq(lulcGewata)
```

	value	count
[1,]	1	396838
[2,]	2	17301
[3,]	3	943
[4,]	4	13645
[5,]	5	470859
[6,]	6	104616
[7,]	NA	817794

```
> hist(lulcGewata)
```

This is a raster with integer values between 1 and 6, but for this raster to be meaningful at all, we need a lookup or attribute table to identify these classes. A .csv file has also been provided as part of the package. Read it in as a data.frame:

```
> data(LUTGewata)
> LUTGewata
```

	ID	Class
1	1	cropland
2	2	bamboo
3	3	bare soil
4	4	coffee plantation
5	5	forest
6	6	wetland

This data.frame represents a lookup table for the raster we just loaded. The ID column corresponds to the values taken on by the lulc raster, and the 'Class' column describes the LULC classes assigned. In R it is possible to add a attribute table to a raster. In order to do this, we need to coerce the raster values to a factor from an integer.

```
> lulc <- as.factor(lulcGewata)
```

If you display the attributes of this raster (just type 'lulc'), it will do so, but will also return an error. This error arises because R expects that a raster with factor values should also have a raster attribute table.

```
> # assign a raster attribute table (RAT)
> levels(lulc) <- LUTGewata
> lulc
```

```
class      : RasterLayer
dimensions : 1177, 1548, 1821996  (nrow, ncol, ncell)
resolution : 30, 30  (x, y)
```

```

extent      : 808755, 855195, 817635, 852945 (xmin, xmax, ymin, ymax)
coord. ref. : +proj=utm +zone=36 +datum=WGS84 +units=m +no_defs +ellps=WGS84 +towgs84=0,0,
data source : in memory
names       : LULC2011_Gewata
values      : 1, 6 (min, max)
attributes  :
      ID    Class
from:  1  cropland
to   :  6  wetland

```

5 Working with multilayered raster data

When working with multispectral or multitemporal raster data, it is convenient to represent multiple raster layers as a single object in R. R works with two types of multilayer raster objects: stacks and bricks. The main difference between the two is that raster stacks can be read from several different data sources (files) and bricks are read from a single file (e.g. a multiband GeoTIFF).

A raster brick from a small area within the Kafa Biosphere Reserve in Southern Ethiopia can be found in the *rasta* package. Set the working directory to the packages home folder and load the raster brick from file by

```

> data(tura)
> # inspect the data
> class(tura) # the object's class

[1] "RasterBrick"
attr(,"package")
[1] "raster"

> projection(tura) # the projection

[1] "+proj=utm +zone=36 +ellps=WGS84 +units=m +no_defs"

> res(tura) # the spatial resolution (x, y)

[1] 30 30

> extent(tura) # the extent of the raster brick

class      : Extent
xmin       : 819105
xmax       : 823395
ymin       : 827745
ymax       : 832185

```

Extracting scene information

This RasterBrick was read from a .grd file. One advantage of this file format (over the GeoTIFF format, for example) is the fact that the specific names of the raster layers making up this brick have been preserved, a feature which is important for identifying raster layers, especially when doing time series analysis (where you need to know the values on the time axis). This RasterBrick was prepared from a Landsat 7 ETM+ time series, and the original scene names were inserted as layer names.

```
> names(tura) # displays the names of all layers in the tura RasterBrick
```

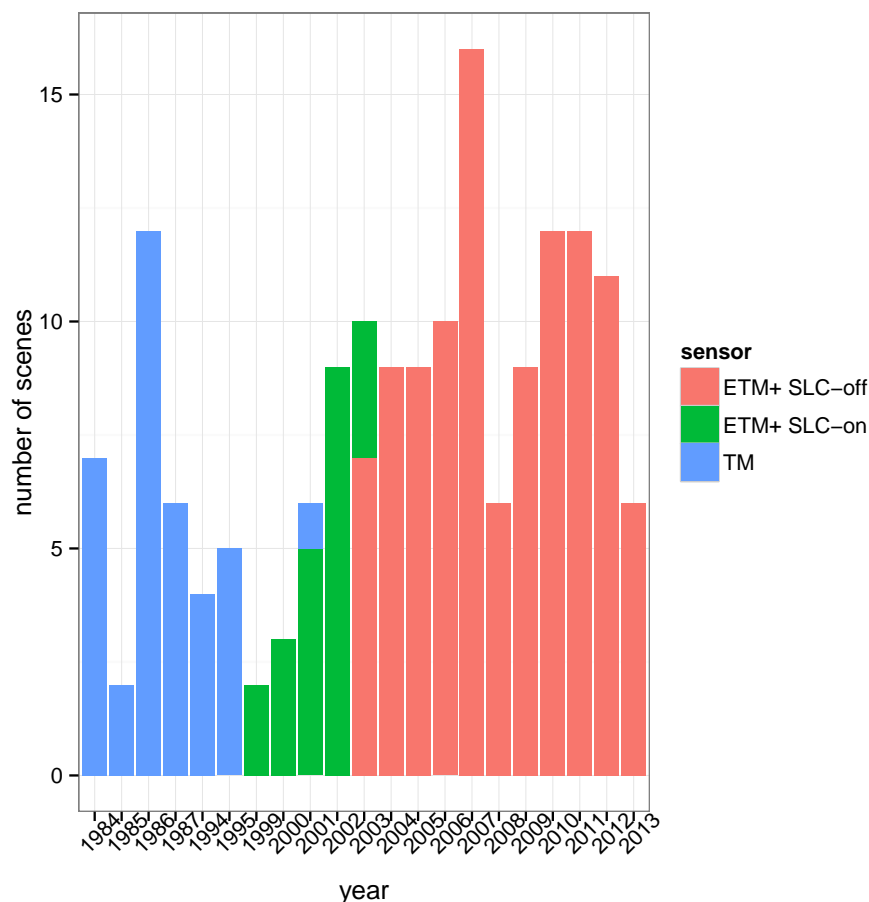
We can parse these names to extract information from them. The first 3 characters indicate which sensor the data come from, with 'LE7' indicating Landsat 7 ETM+ and 'LT5' or 'LT4' indicating Landsat 5 and Landsat 4 TM, respectively. The following 6 characters indicate the path and row (3 digits each), according to the WGS system. The following 7 digits represent the date. The date is formatted in such a way that it equals the year + the julian day. For example, February 5th 2001, aka the 36th day of 2001, would be '2001036'.

```
> # display the 1st 3 characters of the layer names
> sensor <- substr(names(tura), 1, 3)
> print(sensor)
> # display the path and row as numeric vectors in the form (path,row)
> path <- as.numeric(substr(names(tura), 4, 6))
> row <- as.numeric(substr(names(tura), 7, 9))
> print(paste(path, row, sep = ","))
> # display the date
> dates <- substr(names(tura), 10, 16)
> print(dates)
> # format the date in the format yyyy-mm-dd
> as.Date(dates, format = "%Y%j")
```

There is a function in the *rasta* package, `getSceneinfo()` that will parse these names and output a data.frame with all of these attributes.

```
> ?getSceneinfo
> sceneinfo <- getSceneinfo(names(tura))
> print(sceneinfo)
> # add a 'year' column to the sceneinfo dataframe and plot #scenes/year
> sceneinfo$year <- factor(substr(sceneinfo$date, 1, 4), levels = c(1984:2013))

> # barplot with number of scenes per year
> library(ggplot2)
> ggplot(data = sceneinfo, aes(x = year, fill = sensor)) +
+   geom_bar() +
+   labs(y = "number of scenes") +
+   theme_bw() +
+   theme(axis.text.x = element_text(angle = 45))
```



Note that the values along the x-axis of this plot are evenly distributed, even though there are gaps in the values (e.g. between 1987 and 1994). The spacing is due to the fact that we defined `sceneinfo$year` as a vector of *factors* rather than a numeric vectors. Factors act as thematic classes and can be represented by numbers or letters. In this case, the actual values of the factors are not recognized by R. Instead, the levels defined in the `factor()` function define the hierarchy of the factors (in this case we have defined the levels from 1984 up to 2013, according to the range of acquisition dates). For more information on factors in R, check out <http://www.stat.berkeley.edu/classes/s133/factors.html>.

Try to generate the plot above with the years (x-axis) represented as a numeric vector instead of as a factor. Hint: it is not as straightforward as you might think - to convert a factor `x` to a numeric vector, try

```
x <- as.numeric(as.character(x))
```

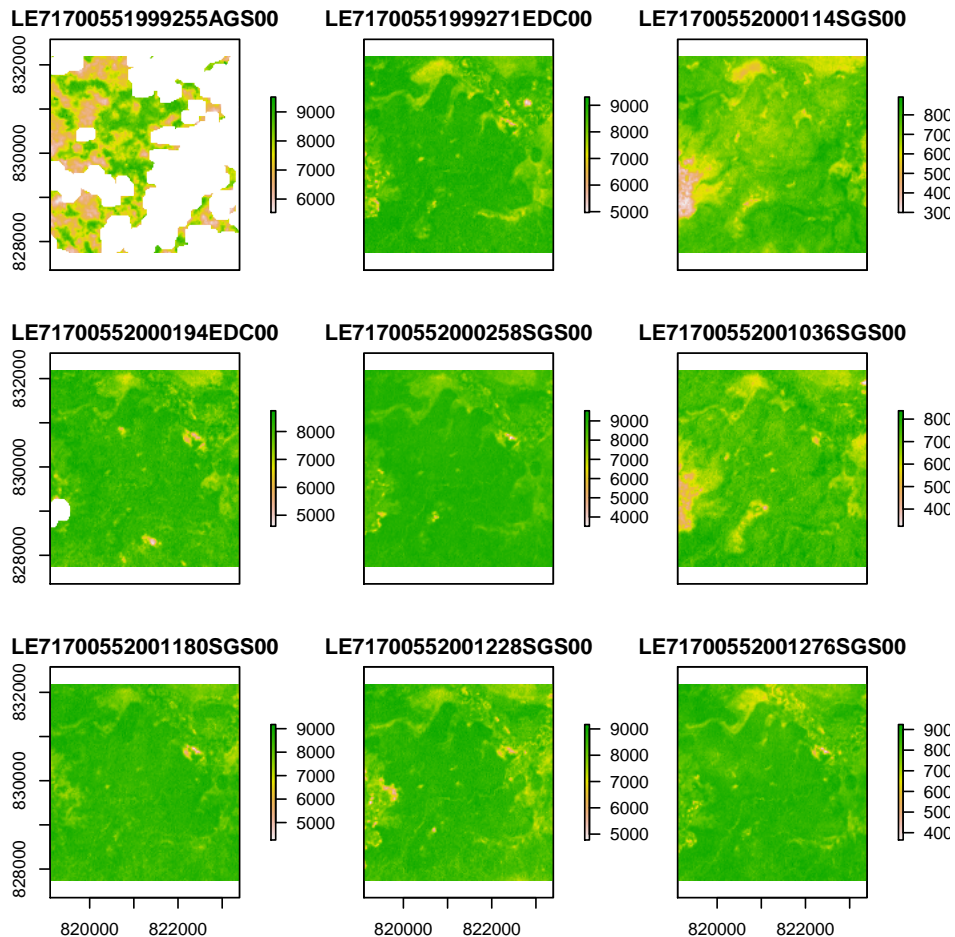
Plotting RasterBricks

A `RasterBrick` can be plotted just as a `RasterLayer`, and the graphics device will automatically split into panels to accommodate the layers (to an extent: R will not attempt to plot 100 layers at once!). To plot the first 9 layers:

```
> plot(tura, c(1:9))
> # alternatively, you can use [[]] notation to specify layers
```



```
> plot(tura[[1:9]])
> # use the information from sceneinfo data.frame to clean up the titles
> plot(tura[[1:9]], main = sceneinfo$date[c(1:9)])
```



Unfortunately, the scale is different for each of the layers, making it impossible to make any meaningful comparison between the raster layers. This problem can be solved by specifying a `breaks` argument in the `plot()` function.

```
> # we need to define the breaks to harmonize the scales (to make the plots comparable)
> bks <- seq(0, 10000, by = 2000) # (arbitrarily) define the breaks
> # we also need to redefine the colour palette to match the breaks
> cols <- rev(terrain.colors(length(bks))) # col = rev(terrain.colors(255)) is the default
> # (opt: check out the RColorBrewer package for other colour palettes)
> # plot again with the new parameters
> plot(tura[[1:9]], main = sceneinfo$date[1:9], breaks = bks, col = cols)
```

Alternatively, the `rasterVis` package has some enhanced plotting functionality for raster objects, including the `levelplot()` function, which automatically provides a common scale for the layers.

```
> library(rasterVis)
> levelplot(tura[[1:6]])
```

```

> # NOTE:
> # for rasterVis plots we must use the [[]] notation for extracting layers

> # providing titles to the layers is done using the 'names.attr' argument in place of 'ma
> levelplot(tura[[1:8]], names.attr = sceneinfo$date[1:8])
> # define a more logical colour scale
> library(RColorBrewer)
> # this package has a convenient tool for defining colour palettes
> ?brewer.pal
> display.brewer.all()
> cols <- brewer.pal(11, 'PiYG')
> # to change the colour scale in levelplot(), we first have to define a rasterTheme objec
> # see ?rasterTheme() for more info
> rtheme <- rasterTheme(region = cols)
> levelplot(tura[[1:8]], names.attr = sceneinfo$date[1:8], par.settings = rtheme)

```

This plot gives us a common scale which allows us to compare values (and perhaps detect trends) from layer to layer. In the above plot, the layer titles do not look very nice – we will solve that problem a bit later.

The rasterVis package has integrated plot types from other packages with the raster package to allow for enhanced analysis of raster data.

```

> # histograms of the first 6 layers
> histogram(tura[[1:6]])
> # box and whisker plot of the first 9 layers
> bwplot(tura[[1:9]])

```

More examples from the rasterVis package can be found @ <http://oscarperpinan.github.io/rastervis/>

Calculating data loss

In this RasterBrick, the layers have all been individually preprocessed from the raw data format into NDVI values. Part of this process was to remove all pixels obscured by clouds or SLC-off gaps (for any ETM+ data acquired after March 2003). For this reason, it may be useful to know how much of the data has been lost to cloud cover and SLC gaps. First, we will calculate the percentage of no-data pixels in each of the layers using the `freq()` function. `freq()` returns a table (matrix) of counts for each value in the raster layer. It may be easier to represent this as a data.frame to access column values.

```

> # try for one layer first
> y <- freq(tura[[1]]) # this is a matrix
> y <- as.data.frame(y)
> # how many NA's are there in this table?
> y$count[is.na(y$value)]

[1] 11340

> # alternatively, using the with() function:
> with(y, count[is.na(value)])

```

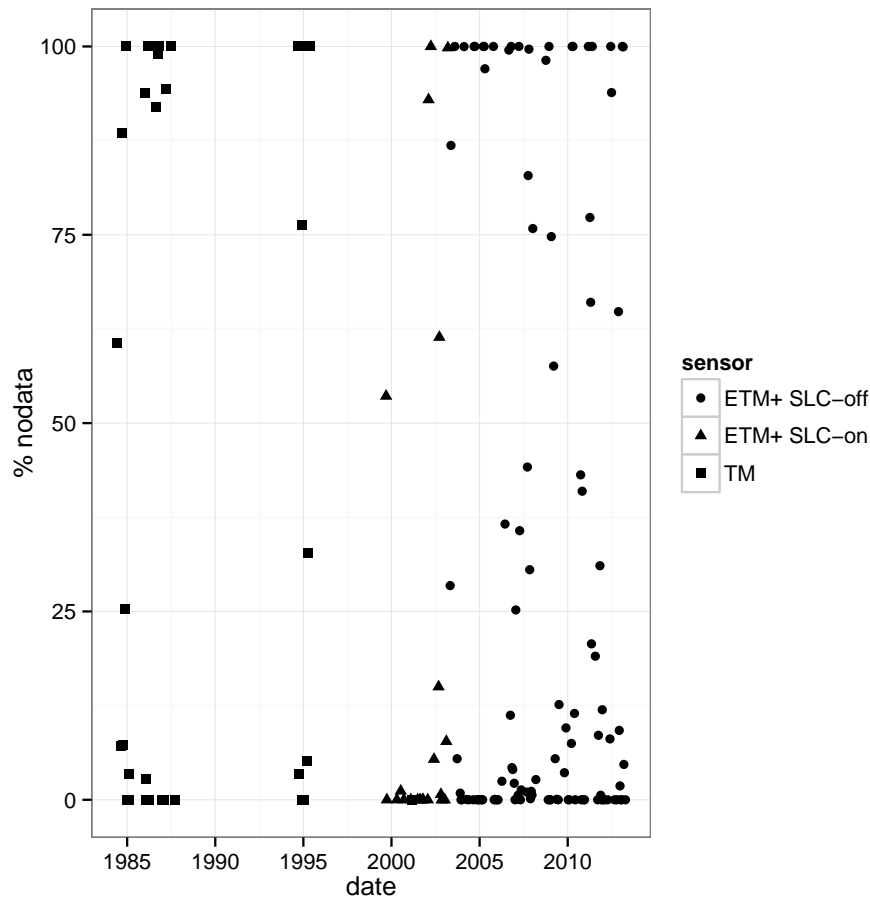
```
[1] 11340
```

```
> # as a %  
> with(y, count[is.na(value)]) / ncell(tura[[1]]) * 100
```

```
[1] 53.58155
```

```
> # apply this over all layers in the RasterBrick  
> # first, prepare a numeric vector to be 'filled' in  
> nas <- vector(mode = 'numeric', length = nlayers(tura))  
> for(i in 1:nlayers(tura)){  
+   y <- as.data.frame(freq(tura[[i]]))  
+   # if there are no NAs, then simply assign a zero  
+   # otherwise, grab the # of NAs from the frequency table  
+   if(!TRUE %in% is.na(y$value)){  
+     nas[i] <- 0  
+   } else {  
+     nas[i] <- with(y, count[is.na(value)]) / ncell(tura[[i]]) * 100  
+   }  
+ }  
> # add this vector as a column in the sceneinfo data.frame (rounded to 2 decimal places)  
> sceneinfo$nodata <- round(nas, 2)
```

```
> # plot these values  
> ggplot(data = sceneinfo, aes(x = date, y = nodata, shape = sensor)) +  
+   geom_point(size = 2) +  
+   labs(y = "% nodata") +  
+   theme_bw()
```



We have now derived some highly valuable information about our time series. For example, we may want to select an image from our time series with relatively little cloud cover to perform a classification. For further time series analysis, the layers with 100% data loss will be of no use to us, so it may make sense to get rid of these layers.

```
> # which layers have 100% data loss?
> which(sceneinfo$nodata == 100)

[1] 13 25 34 35 41 42 44 54 62 78 92 93 101 105 117 126 135 143 147
[20] 149 155 157 162 165

> # supply these indices to the dropLayer() command to get rid of these layers
> tura <- dropLayer(tura, which(sceneinfo$nodata == 100))
> # redefine our sceneinfo data.frame as well
> sceneinfo <- sceneinfo[which(sceneinfo$nodata != 100), ]
> # optional: remake the previous ggplots with this new dataframe
```

With some analyses, it may also be desirable to apply a no-data threshold per scene, in which case layer indices would be selected by:

```
> which(sceneinfo$nodata > some_threshold)
```

In some cases, there may be parts of the study area with more significant data loss due to persistent cloud cover or higher incidence of SLC-off gaps. To map the spatial distribution of data loss, we need to calculate the % of NA in the time series for each *pixel* (ie. looking 'through' the pixel along the time axis). To do this, it is convenient to use the `calc()` function and supply a special function which will count the number of NA's for each pixel along the time axis, divide it by the total number of data in the pixel time series, and output a percentage. `calc()` will output a raster with a percentage no-data value for each pixel.

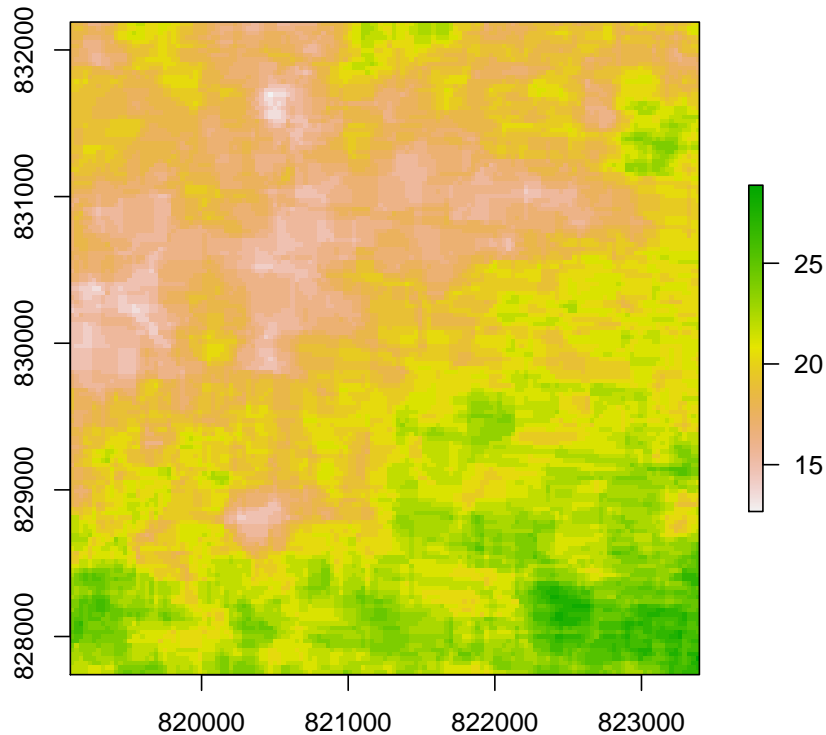
```
> # calc() will apply a function over each pixel
> # in this case, each pixel represents a time series of NDVI values
> # e.g. all values of the 53rd pixel in the raster grid:
> y <- as.numeric(tura[53])
> # how many of these values have been masked (NA)?
> length(y[is.na(y)])
```

```
[1] 23
```

```
> # as a %
> length(y[is.na(y)]) / length(y) * 100
```

```
[1] 16.19718
```

```
> # now wrap this in a calc() to apply over all pixels of the RasterBrick
> nodata <- calc(tura, fun = function(x) length(x[is.na(x)]) / length(x) * 100)
```



Summary statistics

In addition to calculating the % of NA in the time series, there are many other ways we can describe this raster time series. For example, we may want to know the mean or median NDVI value over the whole time series, for example.

...to do....

6 Times Series Analysis

- making time series plots for selected pixels (interactively or otherwise)
- deriving time series statistics
 - linear regression from pixel time series
 - deriving seasonality parameters?
 - creating figures maps using these statistics

There is an easy way to interactively extract data from a single pixel using the raster package. First, a raster plot should be made from which to identify the pixel of interest. Then, the `click()` function allows for extraction of the data contained within that pixel. Simply

calling `click()` with no further arguments will only return the (x, y) coordinates of that point as a 1-row matrix.

```
> plot(tura, 101)
> click()
      x      y
[1,] 819695.2 830544.3
```

We can extract more meaningful information by passing the name of the object (in this case, a `RasterBrick`) as the first argument, followed by the argument `'n = 1'` (or the number of desired points to identify).

```
> plot(tura, 101)
> click(tura, n = 1)
```

In this case, all data from that pixel are extracted. These time series could then easily be coerced (or inserted) into a `data.frame` to plot the data and further analyze them. Let's take a look at the time series of a few different pixels. Instead of using the interactive `click()` function to extract these data, suppose that we know the (x, y) coordinates of a land cover type (or other points of interest, from a field campaign or ground truth dataset, for example).

```
> # several pixel coordinate pairs expressed as separate 1-row matrices
> forest <- matrix(data = c(819935, 832004), nrow = 1, dimnames = list(NULL, c('x', 'y')))
> cropland <- matrix(data = c(819440, 829346), nrow = 1, dimnames = list(NULL, c('x', 'y')))
> wetland <- matrix(data = c(822432, 832076), nrow = 1, dimnames = list(NULL, c('x', 'y')))
> # recall that we can extract pixel data if we know the cell #
> # we can easily convert from xy matrix to cell number with cellFromXY()
> cellFromXY(tura, forest)
> tura[cellFromXY(tura, forest)] # returns a 1-row matrix with all ts values
```

Now we are able to extract the time series data given a set of (x, y) coordinates. Let's put the data from these three points into a `data.frame` to facilitate plotting of the data.

```
> # prepare the data.frame
> ts <- data.frame(sensor = getSceneinfo(names(tura))$sensor,
+                 date = getSceneinfo(names(tura))$date,
+                 forest = t(tura[cellFromXY(tura, forest)]),
+                 cropland = t(tura[cellFromXY(tura, cropland)]),
+                 wetland = t(tura[cellFromXY(tura, wetland)])
+                 )
> print(ts)
> # simple plot of forest time series
> plot(ts$date, ts$forest)
> # same thing but using with()
> with(ts, plot(date, forest))
```

Note the two large gaps in the time series during the 1990's, during which time there are no Landsat data available from the USGS. While we could still use these data to understand historical trends, we will only look at time series data from the ETM+ sensor (ie. data acquired after 1999) for the following exercises.

```
> # remove all data from the TM sensor and plot again
> ts <- ts[which(ts$sensor != "TM"), ]
> with(ts, plot(date, forest))
```

A more informative plot would show these time series side by side with the same scale or on the same plot. These are possible with either the base `plot()` function or using `ggplot2`. Either way, there is some preparation needed, and in the case of `ggplot2`, this may not be immediately obvious. In the following example, we are going to make a `facet_wrap` plot. In order to do so, we need to merge the time series columns to make a `data.frame` with many rows indeed. An additional column will be used to identify the class (forest, cropland or wetland) of each data point, and this class will be used to 'split' the data into 3 facets. The `reshape` package has a convenient function, `melt()`, which will 'automatically' reshape the `data.frame` to make it passable to the `ggplot` framework.

```
> library(reshape)
> # convert dates to characters, otherwise melt() returns an error
> ts$date <- as.character(ts$date)
> tsmelt <- melt(ts)
> head(tsmelt)
```

	sensor	date	variable	value
1	ETM+ SLC-on	1999-09-12	forest	6532
2	ETM+ SLC-on	1999-09-28	forest	8812
3	ETM+ SLC-on	2000-04-23	forest	7564
4	ETM+ SLC-on	2000-07-12	forest	8095
5	ETM+ SLC-on	2000-09-14	forest	8677
6	ETM+ SLC-on	2001-02-05	forest	7279

```
> names(tsmelt) <- c('sensor', 'date', 'class', 'value')
> # convert tsplot$date back to Date class to enable formatting of the plot
> tsmelt$date <- as.Date(tsmelt$date)
> tsplot <- ggplot(data = tsmelt, aes(x = date, y = value / 10000)) +
+   geom_point() +
+   scale_x_date() +
+   labs(y = "NDVI") +
+   facet_wrap(~ class, nrow = 3) +
+   theme_bw()
> tsplot
```

7 Exercise

- to be doable in 3 hours....
- combine concepts from previous lessons as well
- example 1: produce a figure with maximum/minimum/median/mean NDVI per year; figure should have a common scale and be properly labelled

- example 2: compute a time series metric over an entire RasterBrick using the `calc()` function